

coursework_01

January 29, 2026

1 Coursework 1: Image filtering

The coursework includes coding questions and/or written questions. Please read both the text and the code in this notebook to get an idea what you are expected to implement.

1.1 What is expected?

- Complete and run the code using `jupyter-lab`.
- Export (File | Save and Export Notebook As...) the notebook as a PDF file, which contains your code, results and answers, and upload the PDF file onto [Scientia](#).
- Instead of clicking the Export button, you can also run the following command instead:
`jupyter nbconvert coursework_01.ipynb --to pdf`
- If Jupyter complains issues during exporting, it is likely that `pandoc` or `latex` is not installed, or their paths have not been included. You can install the relevant libraries and retry. Alternatively, use the Print function of your browser to export the PDF file.
- If Jupyter-lab does not work for you at the end, alternatively, you can use Google Colab to write the code and export the PDF file.

1.2 Dependencies:

You may need to install [Jupyter-Lab](#) and other libraries used in this coursework, such as by running the command: `pip3 install [package_name]`

```
[1]: # Import libraries (provided)
import imageio.v3 as imageio
import numpy as np
import matplotlib.pyplot as plt
import scipy
import scipy.signal
import math
import time
```

1.3 Q1. Moving average filtering (20 points).

Read the provided clean image, add noise to the image and design a moving average filter for denoising.

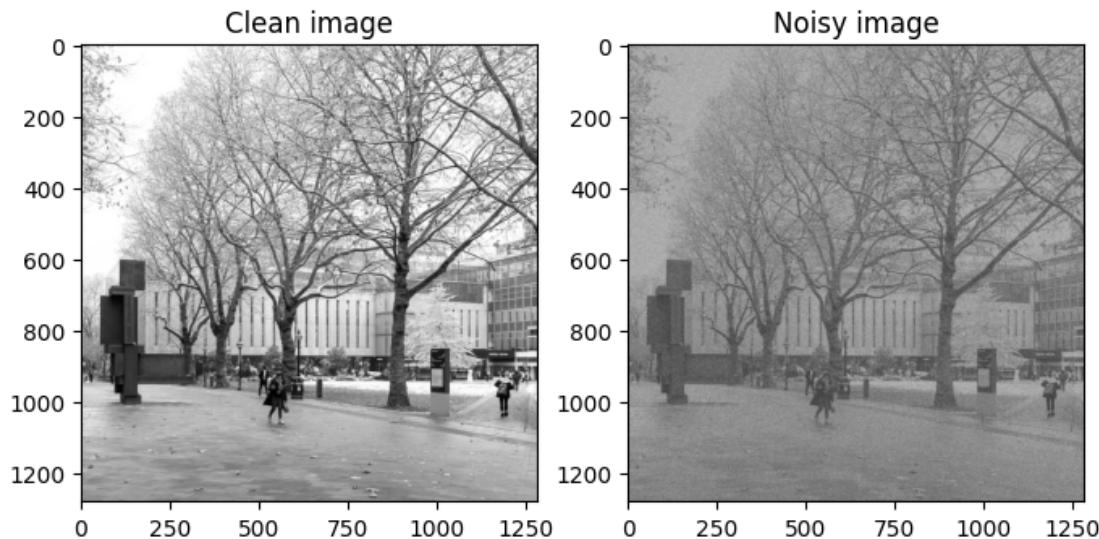
You are expected to design the kernel of the filter and then perform 2D image filtering using the function `scipy.signal.convolve2d()`.

```
[2]: # Read the image (provided)
image = imageio.imread('campus_snow.jpg')

# Corrupt the image with Gaussian noise (provided)
noise_mu = 0
noise_sigma = 50
noise = np.random.normal(noise_mu, noise_sigma, image.shape)
image_noisy = image + noise

# Visualise the images (provided)
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Clean image')

plt.subplot(1, 2, 2)
plt.imshow(image_noisy, cmap='gray')
plt.title('Noisy image')
plt.gcf().set_size_inches(8, 4)
```



1.3.1 Q1.1 Filter the noisy image using a 5x5 moving average filter. Display the filtered image.

```
[3]: # Design the filter h
### Insert your code ####
h = np.ones((5, 5), dtype=np.float64) / 25.0
```

```

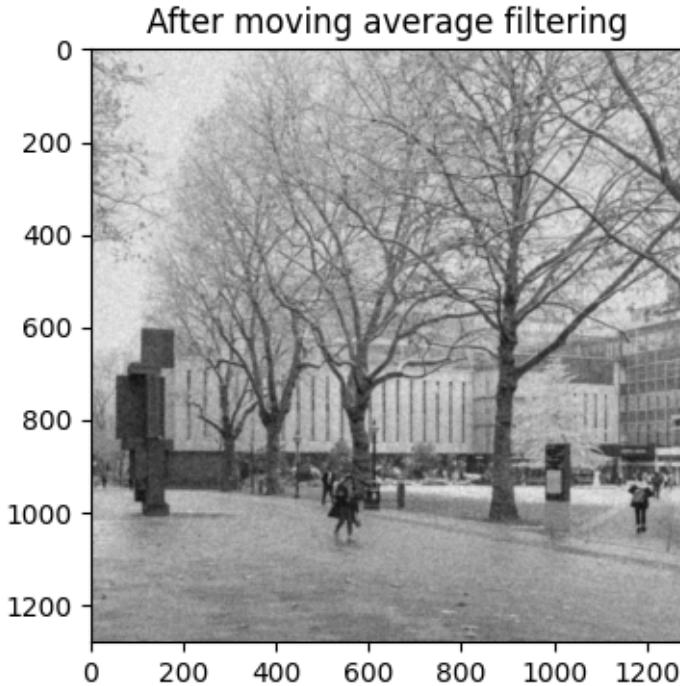
# Convolve the noisy image with h using scipy.signal.convolve2d function
### Insert your code ####
image_filtered = scipy.signal.convolve2d(image_noisy, h, mode='same', boundary='symm')

# Print the filter (provided)
print('Filter h = {0}'.format(h))

# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('After moving average filtering')
plt.gcf().set_size_inches(4, 4)

```

Filter h = [[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]]



1.3.2 Q1.2 Assess the quality of the filtered image using a quantitative metric, the peak signal-to-noise ratio (PSNR).

For this case, the pixel intensity of the image is represented using the uint8 format, with the peak value to be 255. The PSNR is defined as,

$$\text{PSNR} = 10 \cdot \log_{10} \frac{255^2}{\frac{1}{N} \sum_x [J(x) - I(x)]^2}$$

where x denotes the pixel index, N denotes the total number of pixels in the image, J denotes the filtered i.e. denoised image and I denotes the ground truth clean image. The denominator of the term within the logarithm operator is the mean squared error between I and J .

You can find more detail about PSNR [here](#).

```
[4]: # Implement the PSNR function
def eval_psnr(I, J):
    # I: the ground truth clean image (peak value: 255 for uint8 data format)
    # J: the denoised image
    #
    # return: the PSNR metric (unit: dB)

    ### Insert your code ####
    I = I.astype(np.float64)
    J = J.astype(np.float64)
    mse = np.mean((J - I) ** 2)
    if mse == 0:
        return float('inf')
    psnr = 10.0 * np.log10((255.0 ** 2) / mse)
    return psnr

# Evaluate the PSNR for the filtered image (provided)
psnr = eval_psnr(image, image_filtered)

# Print the PSNR (provided)
print('PSNR = {:.2f} dB'.format(psnr))
```

PSNR = 18.48 dB

1.4 Q2. Gaussian filtering (70 points).

1.4.1 Q2.1 Implement a function that constructs a 2D Gaussian filter given the parameter σ .

```
[5]: # Implement the Gaussian filter
def gaussian_filter_2d(sigma):
    # sigma: the parameter sigma for the Gaussian kernel (unit: pixel)
    #
    # return: a 2D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))
    sz = 2 * rad + 1

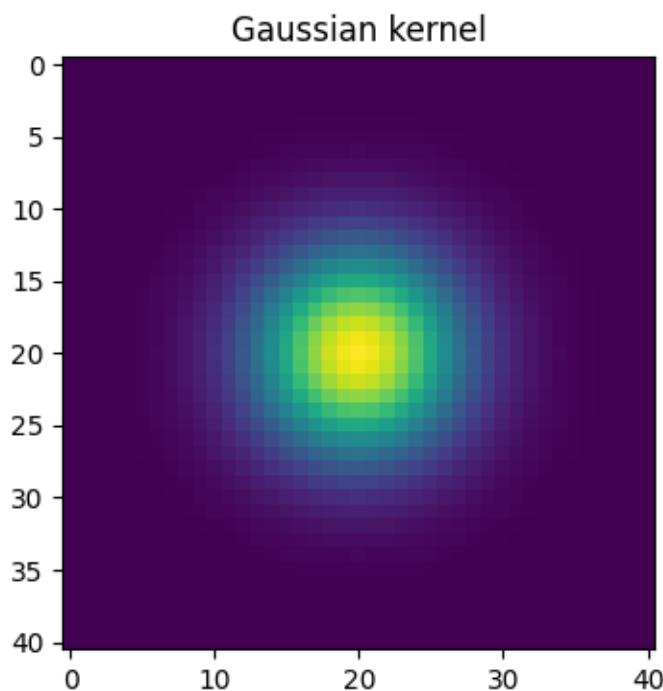
    # Calculate the filter weights
    ### Insert your code ####
```

```

x = np.arange(-rad, rad + 1)
xx, yy = np.meshgrid(x, x)
h = np.exp(-(xx ** 2 + yy ** 2) / (2 * sigma ** 2))
h = h / np.sum(h)
return h

# Visualise the Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)
plt.imshow(h)
plt.title('Gaussian kernel')
plt.gcf().set_size_inches(4, 4)

```



1.4.2 Q2.2 Perform Gaussian filtering ($\sigma = 5$ pixels) for the noisy image, evaluate the computational time for Gaussian filtering and display the filtered image.

```
[6]: # Construct the Gaussian filter (provided)
sigma = 5
h = gaussian_filter_2d(sigma)

# Perform Gaussian filtering and count time
### Insert your code ####
start_time = time.time()
```

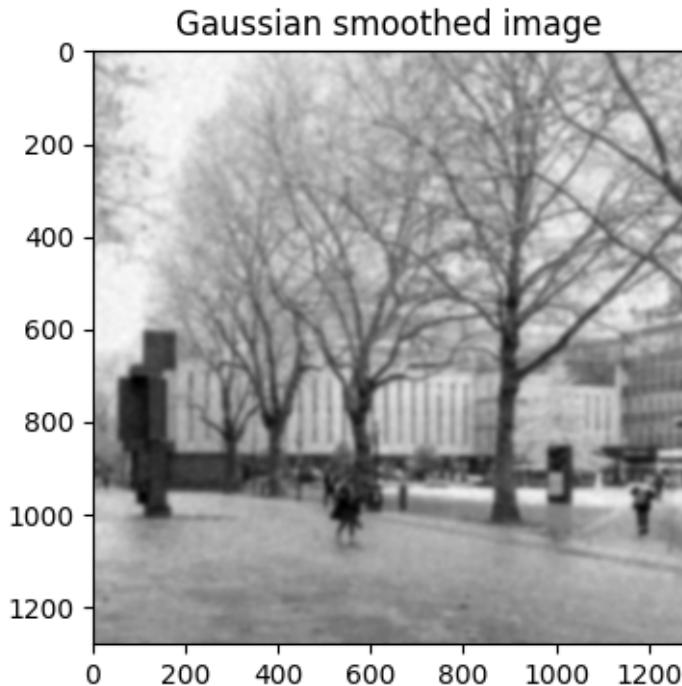
```

image_smoothed_2d = scipy.signal.convolve2d(image_noisy, h, mode='same', boundary='symm')
elapsed_time = time.time() - start_time
print('2D Gaussian filtering time: {:.6f} s'.format(elapsed_time))

# Visualise the filtered image (provided)
plt.imshow(image_smoothed_2d, cmap='gray')
plt.title('Gaussian smoothed image')
plt.gcf().set_size_inches(4, 4)

```

2D Gaussian filtering time: 4.533612 s



1.4.3 Q2.3 Implement a function that generates a 1D Gaussian filter given the parameter σ . Construct 1D Gaussian filters along x-axis and y-axis respectively.

```
[7]: # Implement the 1D Gaussian filter
def gaussian_filter_1d(sigma):
    # sigma: the parameter sigma in the Gaussian kernel (unit: pixel)
    #
    # return: a 1D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))
    sz = 2 * rad + 1
```

```

# Calculate the filter weights
### Insert your code ####
x = np.arange(-rad, rad + 1)
h = np.exp(-(x ** 2) / (2 * sigma ** 2))
h = h / np.sum(h)
return h

# sigma = 5 pixel (provided)
sigma = 5

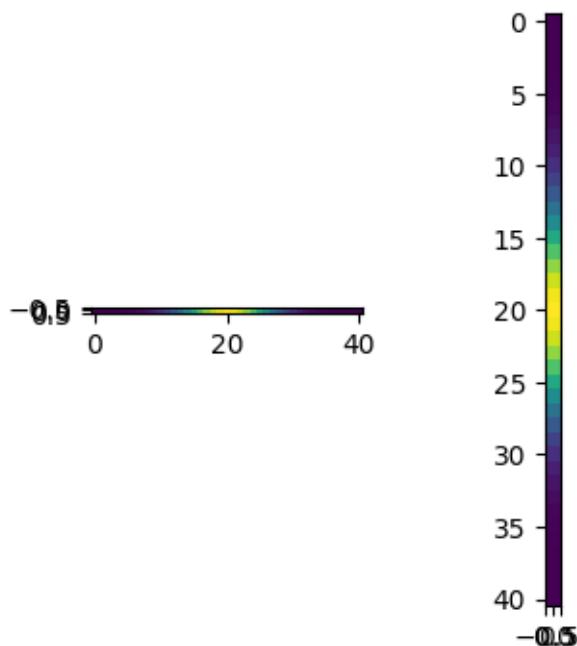
# Construct the Gaussian filter along x-axis. Its shape is (1, sz).
### Insert your code ####
h_x = gaussian_filter_1d(sigma).reshape(1, -1)

# Construct the Gaussian filter along y-axis. Its shape is (sz, 1).
### Insert your code ####
h_y = gaussian_filter_1d(sigma).reshape(-1, 1)

# Visualise the filters (provided)
plt.subplot(1, 2, 1)
plt.imshow(h_x)

plt.subplot(1, 2, 2)
plt.imshow(h_y)
plt.gcf().set_size_inches(4, 4)

```



1.4.4 Q2.4 Perform Gaussian filtering ($\sigma = 5$ pixels) using two separable filters and evaluate the computational time for separable Gaussian filtering. Compare the smoothed image using separable filtering to the smoothed image using a single 2D Gaussian filter.

```
[8]: # Perform separable Gaussian smoothing and count time
#### Insert your code ####
start_time = time.time()
tmp = scipy.signal.convolve2d(image_noisy, h_x, mode='same', boundary='symm')
image_smoothed = scipy.signal.convolve2d(tmp, h_y, mode='same', boundary='symm')
elapsed_time = time.time() - start_time
print('Separable Gaussian filtering time: {:.6f} s'.format(elapsed_time))

# Report the difference between the separably filtered image and the image
# filtered by a single 2D Gaussian filter (provided)
diff = image_smoothed - image_smoothed_2d
print('Mean absolute difference = {:.6f}'.format(np.mean(np.abs(diff))))
```

Separable Gaussian filtering time: 0.401878 s
Mean absolute difference = 0.000000

1.4.5 Q2.5 Perform Gaussian smoothing for the same noisy image, assess the quality of the Gaussian smoothed image using PSNR, when different sigma values are used.

```
[9]: # A list of sigma values (provided)
list_sigma = np.arange(0.5, 5, 0.5)

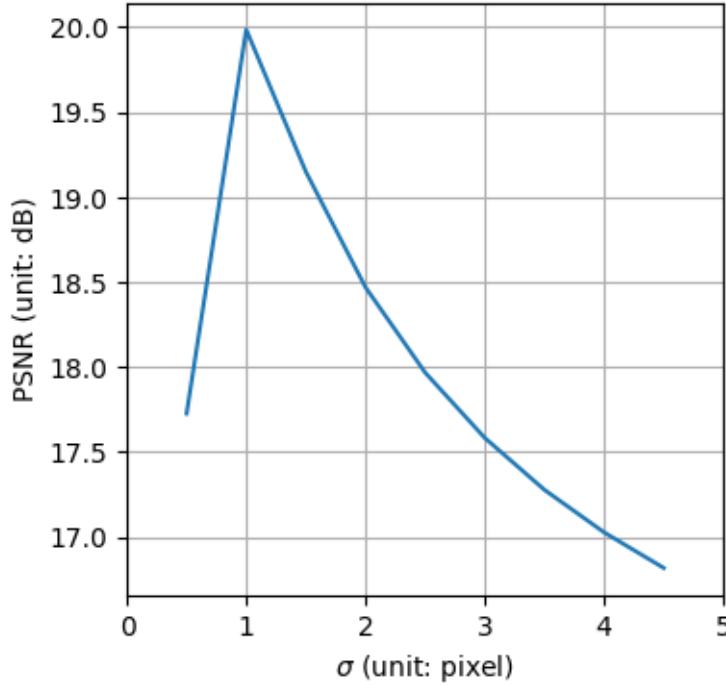
# Perform Gaussian smoothing with different sigma values and record the PSNR
# values
#### Insert your code ####
list_psnr = []
for sigma in list_sigma:
    h = gaussian_filter_2d(sigma)
    image_smoothed = scipy.signal.convolve2d(image_noisy, h, mode='same', boundary='symm')
    list_psnr.append(eval_psnr(image, image_smoothed))

# Plot the PSNR metric against sigma (provided)
plt.plot(list_sigma, list_psnr)
plt.xlim([0, 5])
plt.xlabel('$\sigma$ (unit: pixel)')
plt.ylabel('PSNR (unit: dB)')
plt.grid()
plt.gcf().set_size_inches(4, 4)
```

```

<>:15: SyntaxWarning: invalid escape sequence '\s'
<>:15: SyntaxWarning: invalid escape sequence '\s'
/var/folders/cy/jnf9bq616dbgz88zcmv9rl6w0000gn/T/ipykernel_18141/3062741685.py:1
5: SyntaxWarning: invalid escape sequence '\s'
    plt.xlabel('$\sigma$ (unit: pixel)')

```



1.4.6 Q2.6 Implement 3x3 Sobel filters, perform Sobel filtering for the noisy image, and display the gradient magnitude map.

```

[10]: # Construct the Sobel filters
##### Insert your code #####
sobel_x = np.array([[-1, 0, 1],
                   [-2, 0, 2],
                   [-1, 0, 1]], dtype=np.float64)
sobel_y = np.array([[-1, -2, -1],
                   [ 0,  0,  0],
                   [ 1,  2,  1]], dtype=np.float64)

# Print the filters (provided)
print('Sobel_x = {0}'.format(sobel_x))
print('Sobel_y = {0}'.format(sobel_y))

# Sobel filtering for the noisy image
##### Insert your code #####

```

```

grad_x = scipy.signal.convolve2d(image_noisy, sobel_x, mode='same', boundary='symm')
grad_y = scipy.signal.convolve2d(image_noisy, sobel_y, mode='same', boundary='symm')

# Calculate the gradient magnitude
### Insert your code ####
grad_mag_noisy = np.hypot(grad_x, grad_y)

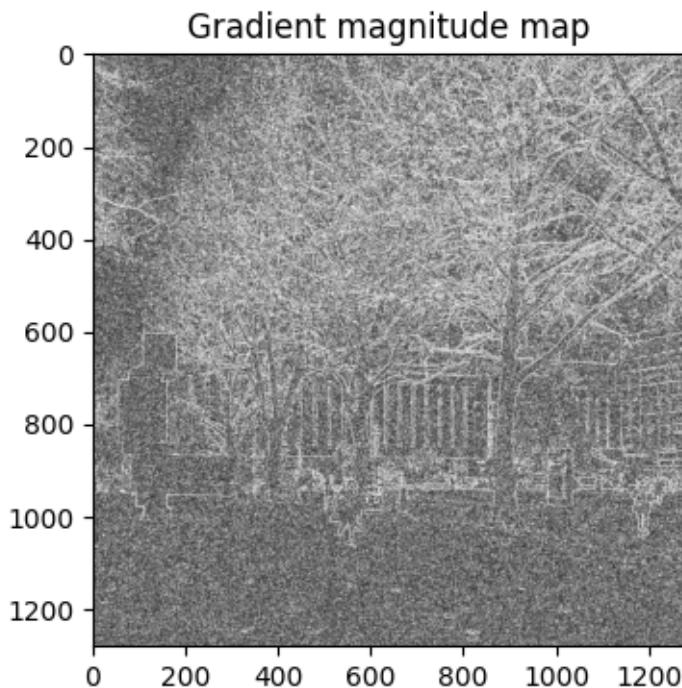
# Display the magnitude map (provided)
plt.imshow(grad_mag_noisy, cmap='gray', vmin=0, vmax=500)
plt.title('Gradient magnitude map')
plt.gcf().set_size_inches(4, 4)

```

```

Sobel_x = [[-1.  0.  1.]
           [-2.  0.  2.]
           [-1.  0.  1.]]
Sobel_y = [[-1. -2. -1.]
           [ 0.  0.  0.]
           [ 1.  2.  1.]]

```



1.4.7 Q2.7 Perform Gaussian smoothing for the noisy image, followed by Sobel filtering and display the gradient magnitude map.

```
[11]: # Parameter for the Gaussian filter (provided)
sigma = 5

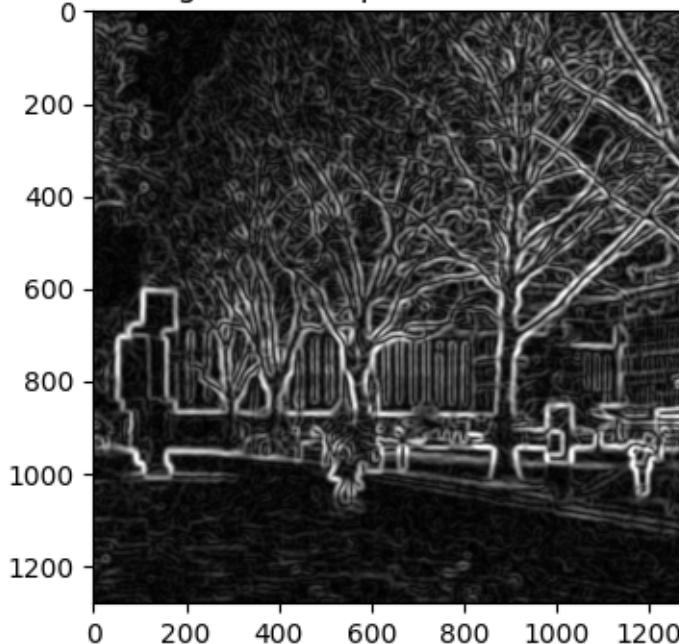
# Gaussian smoothing
### Insert your code ####
h = gaussian_filter_2d(sigma)
image_smoothed = scipy.signal.convolve2d(image_noisy, h, mode='same', boundary='symm')

# Sobel filtering
### Insert your code ####
grad_x = scipy.signal.convolve2d(image_smoothed, sobel_x, mode='same', boundary='symm')
grad_y = scipy.signal.convolve2d(image_smoothed, sobel_y, mode='same', boundary='symm')

# Calculate the gradient magnitude
### Insert your code ####
grad_mag = np.hypot(grad_x, grad_y)

# Display the magnitude map (provided)
plt.imshow(grad_mag, cmap='gray', vmin=0, vmax=100)
plt.title('Gradient magnitude map after Gaussian smoothing')
plt.gcf().set_size_inches(4, 4)
```

Gradient magnitude map after Gaussian smoothing



1.5 Q3. Implement image filters using Pytorch (10 points).

Pytorch is a machine learning framework that supports filtering and convolution.

The Conv2D operator takes an input array of dimension $N \times C_1 \times X \times Y$, applies the filter and outputs an array of dimension $N \times C_2 \times X \times Y$. Here, since we only have one image with one colour channel, we will set $N=1$, $C_1=1$ and $C_2=1$. You can read the documentation of Conv2D for more detail.

```
[12]: # Import libraries (provided)
import torch
```

1.5.1 Q3.1 Expand the dimension of the noisy image into $1 \times 1 \times X \times Y$ and convert it to a Pytorch tensor.

```
[13]: # Expand the dimension of the numpy array
### Insert your code ####
noisy_expanded = np.expand_dims(np.expand_dims(image_noisy, axis=0), axis=0)

# Convert to a Pytorch tensor using torch.from_numpy
### Insert your code ####
tensor_noisy = torch.from_numpy(noisy_expanded.astype(np.float32))
```

1.5.2 Q3.2 Create a Pytorch Conv2D filter, set its kernel to be a 2D Gaussian filter, perform filtering, report computational time and display the result.

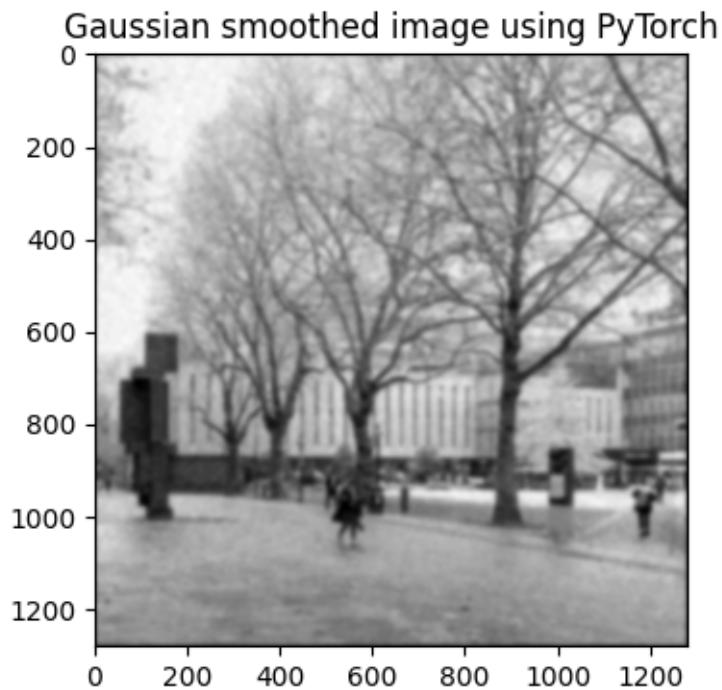
```
[14]: # A 2D Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)

# Construct the Conv2D filter
### Insert your code ####
rad = h.shape[0] // 2
conv = torch.nn.Conv2d(1, 1, kernel_size=h.shape, bias=False, padding=rad)
with torch.no_grad():
    conv.weight[:] = torch.from_numpy(h).float().unsqueeze(0).unsqueeze(0)

# Filtering and assess computational time
### Insert your code ####
start_time = time.time()
with torch.no_grad():
    output = conv(tensor_noisy)
elapsed_time = time.time() - start_time
print('PyTorch Conv2D filtering time: {:.6f} s'.format(elapsed_time))
image_filtered = output.squeeze().cpu().numpy()
```

```
# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('Gaussian smoothed image using PyTorch')
plt.gcf().set_size_inches(4, 4)
```

PyTorch Conv2D filtering time: 135.795004 s



[]: